Blocking: An R Package for Blocking of Records for Record Linkage and Deduplication

by Maciej Beręsewicz and Adam Struzik

Abstract An abstract of less than 250 words.

1 Introduction

Interactive data graphics provides plots that allow users to interact them. One of the most basic types of interaction is through tooltips, where users are provided additional information about elements in the plot by moving the cursor over the plot.

This paper will first review some R packages on interactive graphics and their tooltip implementations. A new package **ToOoOlTiPs** that provides customized tooltips for plot, is introduced. Some example plots will then be given to showcase how these tooltips help users to better read the graphics.

2 Background

Some packages on interactive graphics include **plotly** (Sievert, 2020) that interfaces with Javascript for web-based interactive graphics, **crosstalk** (Cheng and Sievert, 2021) that specializes cross-linking elements across individual graphics. The recent R Journal paper **tsibbletalk** (Wang and Cook, 2021) provides a good example of including interactive graphics into an article for the journal. It has both a set of linked plots, and also an animated gif example, illustrating linking between time series plots and feature summaries.

3 Blocking of records using `blocking` function

4 Integration with existing packages

5 Case study

5.1 Record linkage example

Let us first load the required packages.

```
library(blocking)
library(data.table)
```

We demonstrate the use of `blocking` function for record linkage on the `foreigners` dataset included in the package. This fictional representation of the foreign population in Poland was generated based on publicly available information, preserving the distributions from administrative registers. It contains 110,000 rows with 100,000 entities. Each row represents one record, with the following columns:

- `fname` – first name,
- `sname` – second name,
- `surname` – surname,
- `date` – date of birth,
- `region` – region (county),

- country – country,
- true_id – person ID.

```
data(foreigners)
head(foreigners)

#>      fname  sname     surname       date region country true_id
#>     <char> <char>      <char>     <char> <char>  <char>   <num>
#> 1:    emin            imanov 1998/02/05            031       0
#> 2:  nurlan        suleymanli 2000/08/01            031       1
#> 3:    amio        maharrsmov 1939/03/08            031       2
#> 4:    amik        maharramof 1939/03/08            031       2
#> 5:    amil        maharramov 1993/03/08            031       2
#> 6:   gadir        jahangirov 1991/08/29            031       3
```

We split the dataset into two separate files: one containing the first appearance of each entity in the foreigners dataset, and the other containing its subsequent appearances.

```
foreigners_1 <- foreigners[!duplicated(foreigners$true_id), ]
foreigners_2 <- foreigners[duplicated(foreigners$true_id), ]
```

Now in both datasets we remove slashes from the date column and create a new string column that concatenates the information from all columns (excluding true_id) in each row.

```
foreigners_1[, date := gsub("/", "", date)]
foreigners_1[, txt := paste0(fname, sname, surname, date, region, country)]
foreigners_2[, date := gsub("/", "", date)]
foreigners_2[, txt := paste0(fname, sname, surname, date, region, country)]
head(foreigners_1)

#>      fname  sname     surname     date region country true_id
#>     <char> <char>      <char>   <char> <char>  <char>   <num>
#> 1:    emin            imanov 19980205            031       0
#> 2:  nurlan        suleymanli 20000801            031       1
#> 3:    amio        maharrsmov 19390308            031       2
#> 4:   gadir        jahangirov 19910829            031       3
#> 5:    zaur         bayramova 19961006  01261      031       4
#> 6:    asif           mammadov 19970726            031       5
#>                              txt
#>                           <char>
#> 1:          eminimanov19980205031
#> 2:    nurlansuleymanli20000801031
#> 3:      amiomaharrsmov19390308031
#> 4:     gadirjahangirov19910829031
#> 5: zaurbayramova1996100601261031
#> 6:        asifmammadov19970726031
```

General use

We use the newly created columns in the blocking function, which relies on the default **rnndescent** (Nearest Neighbor Descent) algorithm based on cosine distance. Additionally, we set verbose = 1 to monitor progress. Note that a default parameter of the blocking function is seed = 2023, which sets the random seed.

```
result_reclin <- blocking(x = foreigners_1$txt,
                          y = foreigners_2$txt,
                          verbose = 1)
```

```
#> ===== creating tokens =====
#> ===== starting search (nnd, x, y: 100000, 10000, t: 1232) =====
#> ===== creating graph =====
```

Now we examine the results of record linkage.

- We have created 6,469 blocks.
- The blocking process utilized 1,232 columns (2 character shingles).
- We have 3,916 blocks of 2 elements, 1,604 blocks of 3 elements,..., 2 blocks of 7 elements.

```
result_reclin
```

```
#> ========================================================
#> Blocking based on the nnd method.
#> Number of blocks: 6469.
#> Number of columns used for blocking: 1232.
#> Reduction ratio: 0.9999.
#> ========================================================
#> Distribution of the size of the blocks:
#>    2    3    4    5    6    7
#> 3916 1604  926   19    2    2
```

Structure of the object is as follows:

- `result` – a `data.table` with identifiers and block IDs,
- `method` – name of the ANN algorithm used,
- `deduplication` – whether deduplication was applied,
- `representation` – whether shingles or vectors were used,
- `metrics` – metrics for quality assessment (here `NULL`),
- `confusion` – confusion matrix (here `NULL`),
- `colnames` – column names used for the comparison,
- `graph` – an **igraph** object, mainly for visualization (here `NULL`).

```
str(result_reclin, 1)
```

```
#> List of 8
#> $ result       :Classes 'data.table' and 'data.frame': 10000 obs. of  4 variables:
#>   ..- attr(*, ".internal.selfref")=<externalptr>
#> $ method        : chr "nnd"
#> $ deduplication : logi FALSE
#> $ representation: chr "shingles"
#> $ metrics       : NULL
#> $ confusion     : NULL
#> $ colnames      : chr [1:1232] "0a" "0b" "0c" "0m" ...
#> $ graph         : NULL
#> - attr(*, "class")= chr "blocking"
```

The resulting `data.table` has four columns:

- `x` – reference dataset (i.e. `foreigners_1`) – this may not contain all units of `foreigners_1`,
- `y` – query (each row of `foreigners_2`) – this may not contain all units of `foreigners_2`,
- `block` – block ID,
- `dist` – distance between objects.

```
head(result_reclin$result)
```

```
#>         x     y block      dist
#>     <int> <int> <num>     <num>
#> 1:     3     1     1 0.2216882
#> 2:     3     2     1 0.2122737
#> 3:    21     3     2 0.1172652
#> 4:    57     4     3 0.1863238
#> 5:    57     5     3 0.1379310
#> 6:    61     6     4 0.2307692
```

Let's examine the first pair. Obviously, there are typos in the fname and surname. Nevertheless, the pair is a match.

```
cbind(t(foreigners_1[3, 1:6]), t(foreigners_2[1, 1:6]))
```

```
#>         [,1]        [,2]
#> fname   "amio"      "amik"
#> sname   ""          ""
#> surname "maharrsmov" "maharramof"
#> date    "19390308"  "19390308"
#> region  ""          ""
#> country "031"       "031"
```

Now we use the true_id values to evaluate our approach.

```
matches <- merge(x = foreigners_1[, .(x = 1:.N, true_id)],
                 y = foreigners_2[, .(y = 1:.N, true_id)],
                 by = "true_id")
matches[, block := rleid(x)]
head(matches)
```

```
#> Key: <true_id>
#>    true_id     x     y block
#>      <num> <int> <int> <int>
#> 1:       2     3     1     1
#> 2:       2     3     2     1
#> 3:      20    21     3     2
#> 4:      56    57     4     3
#> 5:      56    57     5     3
#> 6:      60    61     6     4
```

We have 10,000 matched pairs. We use the true_blocks parameter in the blocking function to specify the true block assignments. We obtain the quality metrics for the assessment of record linkage.

```
result_2_reclin <- blocking(x = foreigners_1$txt,
                            y = foreigners_2$txt,
                            verbose = 1,
                            true_blocks = matches[, .(x, y, block)])
```

```
#> ===== creating tokens =====
#> ===== starting search (nnd, x, y: 100000, 10000, t: 1232) =====
#> ===== creating graph =====
```

```
result_2_reclin
```

```
#> ========================================================
```

```
#> Blocking based on the nnd method.
#> Number of blocks: 6469.
#> Number of columns used for blocking: 1232.
#> Reduction ratio: 0.9999.
#> ========================================================
#> Distribution of the size of the blocks:
#>    2    3    4    5    6    7
#> 3916 1604  926   19    2    2
#> ========================================================
#> Evaluation metrics (standard):
#>       recall    precision          fpr          fnr     accuracy specificity
#>      96.7782      78.7000       0.0038       3.2218      99.9957     99.9962
#>     f1_score
#>      86.8079
```

For example, our approach results in a 3.22% false negative rate (FNR). To improve this, we can increase the epsilon parameter of the NND method from 0.1 to 0.5. To do so, we configure the `control_ann` parameter in the `blocking` function using the `controls_ann` and `control_nnd` functions.

```
result_3_reclin <- blocking(x = foreigners_1$txt,
                            y = foreigners_2$txt,
                            verbose = 1,
                            true_blocks = matches[, .(x, y, block)],
                            control_ann = controls_ann(
                              nnd = control_nnd(epsilon = 0.5)))
```

```
#> ===== creating tokens =====
#> ===== starting search (nnd, x, y: 100000, 10000, t: 1232) =====
#> ===== creating graph =====
```

```
result_3_reclin
```

```
#> ========================================================
#> Blocking based on the nnd method.
#> Number of blocks: 6392.
#> Number of columns used for blocking: 1232.
#> Reduction ratio: 0.9999.
#> ========================================================
#> Distribution of the size of the blocks:
#>    2    3    4    5    7
#> 3798 1613  956   21    4
#> ========================================================
#> Evaluation metrics (standard):
#>       recall    precision          fpr          fnr     accuracy specificity
#>      96.8682      80.1100       0.0036       3.1318      99.9960     99.9964
#>     f1_score
#>      87.6957
```

That decreases the FNR to 3.13%.

Integration with the reclin2 package

Let us load the reclin2 package.

```
library(reclin2)
```

Now we present record linkage using the `pair_ann` function. It is based on the `pair_minism` function and reuses some of its source code. The on parameter specifies the column names for the approximate nearest neighbours (ANN) search. Setting deduplication = FALSE enables record linkage. The function works as follows.

```
result_pair_ann <- pair_ann(x = foreigners_1,
                            y = foreigners_2,
                            on = c("fname", "sname", "surname",
                                   "date", "region", "country"),
                            deduplication = FALSE)
head(result_pair_ann)

#>   First data set:  100 000 records
#>   Second data set: 10 000 records
#>   Total number of pairs: 6 pairs
#>   Blocking on: 'fname', 'sname', 'surname', 'date', 'region', 'country'
#>
#>        .x    .y block
#>     <int> <int> <num>
#> 1:      3     1     1
#> 2:      3     2     1
#> 3:     21     3     2
#> 4:     57     4     3
#> 5:     57     5     3
#> 6:     61     6     4
```

The `pair_ann` function returns the total number of pairs. This output can be integrated into the pipeline of the **reclin2** package. We compare pairs across all selected variables using the Jaro-Winkler distance. The similarity scores are summed across the variables and we set threshold = 4.5 to accept a pair.

```
selected_pair_ann <- result_pair_ann |>
  compare_pairs(on = c("fname", "sname", "surname",
                       "date", "region", "country"),
                comparators = list(cmp_jarowinkler())) |>
  score_simple("score",
               on = c("fname", "sname", "surname",
                      "date", "region", "country")) |>
  select_threshold("threshold", score = "score", threshold = 4.5) |>
  link(selection = "threshold")

head(selected_pair_ann)

#>   Total number of pairs: 6 pairs
#>
#> Key: <.y>
#>        .y    .x   fname.x sname.x  surname.x   date.x region.x country.x
#>     <int> <int>    <char>  <char>     <char>   <char>   <char>    <char>
#> 1:      1     3      amio          maharrsmov 19390308                031
#> 2:      2     3      amio          maharrsmov 19390308                031
#> 3:      3    21      amil            khalilov 19990901    01465        031
#> 4:      4    57  javansjir         m kayilov 19691011                031
#> 5:      5    57  javansjir         m kayilov 19691011                031
#> 6:      6    61    rashad           mehtiyev 19980320                031
#>    true_id.x                      txt.x   fname.y sname.y  surname.y
#>        <num>                     <char>    <char>  <char>     <char>
```

```
#> 1:         2     amiomaharrsmov19390308031       amik       maharramof
#> 2:         2     amiomaharrsmov19390308031       amil       maharramov
#> 3:        20   amilkhalilov1999090101465031       amul        khalilpv
#> 4:        56 javansjirm kayilov19691011031    javanshir       mikayilov
#> 5:        56 javansjirm kayilov19691011031    javsnshir       m kayilov
#> 6:        60      rashadmehtiyev19980320031      rasgad        meht9yev
#>      date.y region.y country.y true_id.y                        txt.y
#>      <char>   <char>    <char>     <num>                       <char>
#> 1: 19390308               031         2     amikmaharramof19390308031
#> 2: 19930308               031         2     amilmaharramov19930308031
#> 3: 19990901    01465      031        20   amulkhalilpv1999090101465031
#> 4: 19961011               031        56 javanshirmikayilov19961011031
#> 5: 19691011               031        56 javsnshirm kayilov19691011031
#> 6: 19890320               031        60    rasgadmeht9yev19890320031
```

We observe that the example pairs are matches.

5.2 Deduplication example

We demonstrate deduplication using the blocking function on the RLdata500 dataset from
the **RecordLinkage** package. Note that the dataset is included in the blocking package.
It contains artificial personal data. Fifty records have been duplicated with randomly
generated errors. Each row represents one record, with the following columns:

- fname_c1 – first name, first component,
- fname_c2 – first name, second component,
- lname_c1 – last name, first component,
- lname_c2 – last name, second component,
- by – year of birth,
- bm – month of birth,
- bd – day of birth,
- rec_id – record id,
- ent_id – entity id.

```
data(RLdata500)
head(RLdata500)
```

```
#>     fname_c1 fname_c2 lname_c1 lname_c2    by    bm    bd rec_id ent_id
#>       <char>   <char>   <char>   <char> <int> <int> <int>  <int>  <int>
#> 1:   CARSTEN             MEIER            1949     7    22      1     34
#> 2:      GERD             BAUER            1968     7    27      2     51
#> 3:    ROBERT           HARTMANN           1930     4    30      3    115
#> 4:    STEFAN             WOLFF            1957     9     2      4    189
#> 5:      RALF           KRUEGER            1966     1    13      5     72
#> 6:   JUERGEN            FRANKE            1929     7     4      6    142
```

We create a new column (id_count) that indicates how many times a given unit occurs
and then add leading zeros to the bm and bd columns. Finally, we create a new string column
that concatenates the information from all columns (excluding rec_id, ent_id and id_count)
in each row.

```
RLdata500[, id_count :=.N, ent_id]
RLdata500[, bm:=sprintf("%02d", bm)]
RLdata500[, bd:=sprintf("%02d", bd)]
RLdata500[, txt:=tolower(
  paste0(fname_c1,fname_c2,lname_c1,lname_c2,by,bm,bd))]
head(RLdata500)
```

```
#>    fname_c1 fname_c2 lname_c1 lname_c2     by      bm      bd rec_id ent_id
#>      <char>   <char>   <char>   <char> <int> <char> <char>  <int>  <int>
#> 1:  CARSTEN            MEIER             1949     07     22      1     34
#> 2:    GERD             BAUER            1968     07     27      2     51
#> 3:  ROBERT           HARTMANN           1930     04     30      3    115
#> 4:  STEFAN             WOLFF            1957     09     02      4    189
#> 5:    RALF            KRUEGER           1966     01     13      5     72
#> 6: JUERGEN            FRANKE            1929     07     04      6    142
#>   id_count                     txt
#>      <int>                  <char>
#> 1:        1  carstenmeier19490722
#> 2:        2     gerdbauer19680727
#> 3:        1 roberthartmann19300430
#> 4:        1    stefanwolff19570902
#> 5:        1    ralfkrueger19660113
#> 6:        1  juergenfranke19290704
```

As in the previous example, we use the txt column in the blocking function. This time, we set ann = hnsw to use the Hierarchical Navigable Small World (HNSW) algorithm from the **RcppHNSW** package and graph = TRUE to obtain an **igraph** object for visualization.

```
result_dedup_hnsw <- blocking(x = RLdata500$txt,
                              ann = "hnsw",
                              graph = TRUE,
                              verbose = 1)
```

```
#> ===== creating tokens =====
#> ===== starting search (hnsw, x, y: 500, 500, t: 429) =====
#> ===== creating graph =====
```

The results are as follows.

```
result_dedup_hnsw
```

```
#> ========================================================
#> Blocking based on the hnsw method.
#> Number of blocks: 133.
#> Number of columns used for blocking: 429.
#> Reduction ratio: 0.9916.
#> ========================================================
#> Distribution of the size of the blocks:
#>  2  3  4  5  6  7  8  9 10 11 12 17
#> 46 35 23  8  6  6  2  3  1  1  1  1
```

```
head(result_dedup_hnsw$result)
```

```
#>        x     y block        dist
#>    <int> <int> <num>       <num>
#> 1:     1    64    35 0.47379863
#> 2:     2    43     1 0.08074522
#> 3:     2   486     1 0.41023219
#> 4:     3   450    88 0.43263358
#> 5:     4    50    13 0.52565831
#> 6:     5   128     2 0.51333570
```

Now we visualize connections using the obtained graph.

```
plot(result_dedup_hnsw$graph, vertex.size = 1, vertex.label = NA)
```


Figure 1: Connection graph

We create a long `data.table` with information on blocks and units from the original dataset.

```
df_block_melted <- melt(result_dedup_hnsw$result, id.vars = c("block", "dist"))
df_block_melted_rec_block <- unique(df_block_melted[, .(rec_id=value, block)])
head(df_block_melted_rec_block)

#>    rec_id block
#>     <int> <num>
#> 1:      1    35
#> 2:      2     1
#> 3:      3    88
#> 4:      4    13
#> 5:      5     2
#> 6:      6    35
```

We add the block information to the final dataset.

```
RLdata500[df_block_melted_rec_block, on = "rec_id", block_id := i.block]
head(RLdata500)

#>    fname_c1 fname_c2 lname_c1 lname_c2    by     bm     bd rec_id ent_id
#>      <char>   <char>   <char>   <char> <int> <char> <char>  <int>  <int>
#> 1: CARSTEN            MEIER             1949     07     22      1     34
```

```
#> 2:     GERD          BAUER          1968    07    27    2    51
#> 3:   ROBERT       HARTMANN          1930    04    30    3   115
#> 4:   STEFAN         WOLFF          1957    09    02    4   189
#> 5:     RALF        KRUEGER          1966    01    13    5    72
#> 6:  JUERGEN         FRANKE          1929    07    04    6   142
#>    id_count                   txt block_id
#>       <int>                <char>    <num>
#> 1:        1   carstenmeier19490722       35
#> 2:        2      gerdbauer19680727        1
#> 3:        1 roberthartmann19300430       88
#> 4:        1    stefanwolff19570902       13
#> 5:        1     ralfkrueger19660113        2
#> 6:        1 juergenfranke19290704       35
```

We can check in how many blocks the same entities (ent_id) are observed. In our example, all the same entities are in the same blocks.

```
RLdata500[, .(uniq_blocks = uniqueN(block_id)), .(ent_id)][, .N, uniq_blocks]
```

```
#>    uniq_blocks     N
#>          <int> <int>
#> 1:           1   450
```

Now we can visualize the distances between the units stored in the result_dedup_hnsw$result dataset. Clearly we have a mixture of two groups: matches (close to 0) and non-matches (close to 1).

```
hist(result_dedup_hnsw$result$dist, xlab = "Distances",
     ylab = "Frequency", breaks = "fd",
     main = "Distances calculated between units")
```

## Distances calculated between units

Figure 2: Distances calculated between units

Finally, we visualize the result based on the information whether a block contains matches or not.

```
df_for_density <- copy(df_block_melted[block %in% RLdata500$block_id])
df_for_density[, match:= block %in% RLdata500[id_count == 2]$block_id]

plot(density(df_for_density[match==FALSE]$dist),
    col = "blue", xlim = c(0, 0.8),
    main = "Distribution of distances between\n
    clusters type (match=red, non-match=blue)")
lines(density(df_for_density[match==TRUE]$dist),
     col = "red", xlim = c(0, 0.8))
```
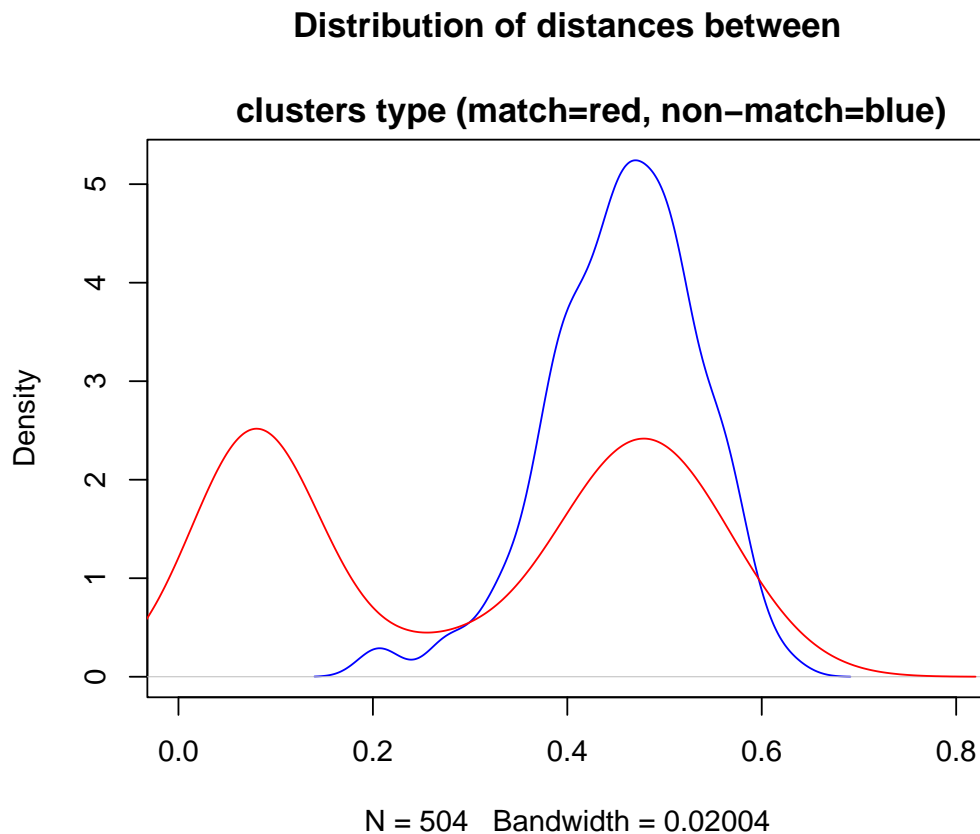
Distribution of distances between

clusters type (match=red, non−match=blue)

N = 504   Bandwidth = 0.02004

Figure 3: Distribution of distances between clusters type

Now we compare the evaluation metrics across all ANN algorithms supported by the blocking function, i.e. NND, HNSW, Approximate Nearest Neighbors Oh Yeah (Annoy, from the **RcppAnnoy** package), Locality-sensitive hashing (LSH, from the **mlpack** package), and k-Nearest Neighbors (kNN – denoted as "kd", from the **mlpack** package). We use the rec_id and ent_id columns from the RLdata500 dataset to specify the true blocks and then calculate evaluation metrics for all algorithms. Additionally, we assess blocking using the klsh function from the **klsh** package, configured to create 10 blocks and 100 blocks, respectively. In both settings, we use 20 random projections and 2-character shingles. The results are as follows (klsh_10 and klsh_100 refer to the klsh algorithm with 10 blocks and 100 blocks, respectively).

```
true_blocks <- RLdata500[, c("rec_id", "ent_id"), with = FALSE]
setnames(true_blocks, old = c("rec_id", "ent_id"), c("x", "block"))
eval_metrics <- list()
ann <- c("nnd", "hnsw", "annoy", "lsh","kd")
for (algorithm in ann) {
  eval_metrics[[algorithm]] <- blocking(x = RLdata500$txt,
                              ann = algorithm,
                              true_blocks = true_blocks)$metrics
}

set.seed(2025)
blocks_klsh_10 <- klsh::klsh(
  r.set = RLdata500[, c("fname_c1", "fname_c2", "lname_c1",
                        "lname_c2", "by", "bm", "bd")],
  p = 20,
  num.blocks = 10,
```

Table 1: A basic table

species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex	year
Adelie	Torgersen	39.1	18.7	181	3750	male	2007
Adelie	Torgersen	39.5	17.4	186	3800	female	2007
Adelie	Torgersen	40.3	18.0	195	3250	female	2007
Adelie	Torgersen	NA	NA	NA	NA	NA	2007
Adelie	Torgersen	36.7	19.3	193	3450	female	2007
Adelie	Torgersen	39.3	20.6	190	3650	male	2007

```
  k = 2)
klsh_10_metrics <- klsh::confusion.from.blocking(
  blocking = blocks_klsh_10,
  true_ids = RLdata500$ent_id)[-1]
klsh_10_metrics$f1_score <- 2 * klsh_10_metrics$precision *
  klsh_10_metrics$recall /
  (klsh_10_metrics$precision + klsh_10_metrics$recall)
eval_metrics$klsh_10 <- unlist(klsh_10_metrics)
blocks_klsh_100 <- klsh::klsh(
  r.set = RLdata500[, c("fname_c1", "fname_c2", "lname_c1",
                        "lname_c2", "by", "bm", "bd")],
  p = 20,
  num.blocks = 100,
  k = 2)
klsh_100_metrics <- klsh::confusion.from.blocking(
  blocking = blocks_klsh_100,
  true_ids = RLdata500$ent_id)[-1]
klsh_100_metrics$f1_score <- 2 * klsh_100_metrics$precision *
  klsh_100_metrics$recall /
  (klsh_100_metrics$precision + klsh_100_metrics$recall)
eval_metrics$klsh_100 <- unlist(klsh_100_metrics)

do.call(rbind, eval_metrics) * 100

#>          recall precision        fpr fnr accuracy specificity f1_score
#> nnd         100 5.0607287 0.7522053   0 99.24810    99.24779 9.633911
#> hnsw        100 4.7573739 0.8027265   0 99.19760    99.19727 9.082652
#> annoy       100 4.8030740 0.7947073   0 99.20561    99.20529 9.165903
#> lsh          98 1.1207685 3.4667201   2 96.53387    96.53328 2.216192
#> kd          100 4.3066322 0.8909383   0 99.10942    99.10906 8.257638
#> klsh_10      82 0.3290794 9.9582999  18 90.03848    90.04170 0.655528
#> klsh_100     86 3.4649476 0.9607057  14 99.03407    99.03929 6.661503
```

6   Customizing tooltip design with ToOoOlTiPs

ToOoOlTiPs is a packages for customizing tooltips in interactive graphics, it features these possibilities.

7   A gallery of tooltips examples

The **palmerpenguins** data (Horst et al., 2020) features three penguin species which has a lovely illustration by Alison Horst in Figure 4.

Table 1 prints at the first few rows of the penguins data:

Figure 5 shows an plot of the penguins data, made using the **ggplot2** package.

```
penguins %>%
```

Figure 4: Artwork by allison_horst

```
ggplot(aes(x = bill_depth_mm, y = bill_length_mm,
           color = species)) +
geom_point()
```

8 Summary

We have displayed various tooltips that are available in the package **ToOoOlTiPs**.

9 Acknowledgements

References

J. Cheng and C. Sievert. *crosstalk: Inter-Widget Interactivity for HTML Widgets*, 2021. URL https://CRAN.R-project.org/package=crosstalk. R package version 1.1.1. [p1]

A. M. Horst, A. P. Hill, and K. B. Gorman. *palmerpenguins: Palmer Archipelago (Antarctica) penguin data*, 2020. URL https://allisonhorst.github.io/palmerpenguins/. R package version 0.1.0. [p13]

C. Sievert. *Interactive Web-Based Data Visualization with R, plotly, and shiny*. Chapman and Hall/CRC, 2020. ISBN 9781138331457. URL https://plotly-r.com. [p1]

E. Wang and D. Cook. Conversations in time: interactive visualisation to explore structured temporal data. *The R Journal*, 2021. doi: 10.32614/RJ-2021-050. URL https://journal.r-project.org/archive/2021/RJ-2021-050/index.html. [p1]

Maciej Beręsewicz
University of Economics and BusinessStatisical Office in Poznań
Department of Statistics, Poznań, Poland
*Centre for the Methodology of Population Studies*

Figure 5: A basic non-interactive plot made with the ggplot2 package on palmer penguin data. Three species of penguins are plotted with bill depth on the x-axis and bill length on the y-axis. Visit the online article to access the interactive version made with the plotly package.

https://maciejberesewicz.com
ORCiD: 0000-0002-8281-4301
maciej.beresewicz@poznan.pl

Adam Struzik
Adam Mickiewicz UniversityStatisical Office in Poznań
Department of Mathematics, Poznań, Poland
Centre for Urban Statistics
adastr5@st.amu.edu.pl